

FE542HW2XYX

February 14, 2018

1 FE542 Time Series Homework 2

Yuxuan Xia

1.1 Problem 1

Basic libraries


```
In [68]: import numpy as np
import pandas as pd
import scipy as sp
import matplotlib.pyplot as plt
import statsmodels.api as sm
```

- (a) Generate 2000 observations from an exponential distribution with mean 10. Report the skewness and kurtosis.

```
In [23]: from statsmodels.stats.stattools import jarque_bera
def jarque_bera_test(data):
    jbres = jarque_bera(obs)
    print("JB statistic:", jbres[0])
    print("JB p-value:", jbres[1])
    print("skew:", jbres[2])
    print("kurtosis:", jbres[3])

def hist_check(data):
    plt.hist(data, normed=True)
    x=np.linspace(-3,3,100)
    y=sp.stats.norm.pdf(x)
    plt.plot(x,y,color='red')
    plt.show()

In [29]: obs_expon = sp.stats.expon.rvs(scale=10,size=2000)
hist_check(obs_expon)
jarque_bera_test(obs_expon)
```




```
JB statistic: 60.46516029297574
JB p-value: 7.41578736531e-14
skew: -0.12724513749466718
kurtosis: 3.812906282083957
```

- (b) Generate 2000 observations from a central t-distribution with 8 degrees of freedom. Test if the distribution of the generated numbers is normal. Use a Jarque-Berra test and a Shapiro Wilk test.

Briefly describe the later test. You do not have to use mathematical formulas only the idea of the test.

```
In [101]: def shapiro_wilk_test(data):
            res = sp.stats.shapiro(obs_t)
            print("Shapiro statistic:", res[0])
            print("Shapiro p-value:", res[1])

In [102]: obs_t = sp.stats.t.rvs(df=8, size=2000)
            hist_check(obs_t)
            jarque_bera_test(obs_t)
            shapiro_wilk_test(obs_t)
```



```
JB statistic: 60.46516029297574
JB p-value: 7.41578736531e-14
skew: -0.12724513749466718
kurtosis: 3.812906282083957
Shapiro statistic: 0.99196857213974
Shapiro p-value: 4.918607299231326e-09
```

From both results of Jarque Bera test and Shapiro Wilk test, the null-hypothesis (normality) is rejected under the chosen level 0.01%. So, the random variables generated from t-distribution with degrees of freedom at 8 are not normal.

Ideal: Small values of the Shapiro test statistic (i.e. W statistic) are evidence of departure from normality and percentage points for the W statistic, obtained via Monte Carlo simulations, were reproduced by Pearson and Hartley. This test has done very well in comparison studies with other goodness of fit tests.

- (c) Test if the previously generated time series are independent. Use a Ljung-Box portmanteau test.

```
In [277]: from statsmodels.tsa.stattools import acf
def ljung_box_q_test(data,m,p=0):
    """p: number of efficients of your model"""
    res = acf(data,nlags=m,qstat=True)
    autocorr = res[0]
    n = len(data)
    Q = n*(n+2)*np.sum([autocorr[k]**2/(n-k) for k in range(1,m+1)])
```

```

print("autocorrelation:",res[0])
print("Ljung Box q statistic:",res[1])
print("one-tailed p-value:",[0.5-abs(0.5-sp.stats.chi2.cdf(val,df=i+1)
                                for i,val in enumerate(res[1]))])

```

```

In [165]: ljung_box_q_test(obs_expon,4,0)
          ljung_box_q_test(obs_t,5,0)

```

```

autocorrelation: [ 1.          0.03889566 -0.01271072 -0.00187121  0.02574125]
Ljung Box q statistic: [ 3.03028541  3.35405694  3.36107734  4.69028463]
one-tailed p-value: [0.081723639788411018, 0.18692861571470631, 0.33923188257387138
autocorrelation: [ 1.          -0.02080509  0.03434603  0.01131843  0.01960427 -0.01
Ljung Box q statistic: [ 0.86700307  3.23102577  3.487881    4.25884636  4.49598007
one-tailed p-value: [0.35178632244481878, 0.19878868825961227, 0.32233731976394697,

```

Since p-value are higher than 1%, the null hypothesis of independency is not rejected

- (d) Create a new time series in the following way. If x_t are the generated numbers in part a create a new time series:

$$y_t = x_t + 0.4x_{t-1} + \epsilon_t$$

; where $\epsilon_t \sim N(0; \sigma^2)$. Test if the new time series y_t has independent terms. To this end vary σ from small (e.g., 0.1) to large. What do you observe? Create a table to clearly display the phenomenon.

```

In [174]: def gen_hybrid_time_series(xt,sigma):
          length = len(xt)-1
          eps = sp.stats.norm.rvs(loc=0,scale=sigma,size=length)
          yt=xt[1:]+0.4*xt[:-1]+eps[:]
          return yt

```

```

In [175]: from statsmodels.tsa.stattools import adfuller

```

```

def stationarity_test(time_series,regression):
    # determine rolling statics
    rolmean = pd.rolling_mean(time_series,window=200)
    rolstd = pd.rolling_std(time_series,window=200)
    # rolstd = pd.rolling(window=12).std()
    # plot rolling statistics
    fig = plt.figure(figsize=(12,8))
    orig = plt.plot(time_series,color='blue',label="original")
    mean = plt.plot(rolmean,color='red',label="rolling mean")
    std = plt.plot(rolstd,color='black',label="rolling std")
    plt.legend(loc='best')
    plt.show()

    # perform Dickey-Fuller test

```

```

dfctest = adfuller(time_series, regression='c', autolag="BIC")
dfoutput = pd.Series(dfctest[0:4], index=['Test Statistic', 'p-value',
for key, value in dfctest[4].items():
    dfoutput['Critical Value (%s)' % key] = value
print(dfoutput)

```

In [176]: # autocorrelation


```

def autocorrelation_check(timeseries):
    fig = plt.figure(figsize=(12,8))
    ax1 = fig.add_subplot(211)
    fig = sm.graphics.tsa.plot_acf(timeseries, lags=40, ax=ax1)
    ax2 = fig.add_subplot(212)
    fig = sm.graphics.tsa.plot_pacf(timeseries, lags=40, ax=ax2)
    plt.show()

```

In [177]: def autocorrelation_independency_unitroot_check(timeseries):
autocorrelation_check(timeseries)
ljung_box_q_test(timeseries,10,0)
stationarity_test(timeseries,'nc')

In [178]: xt = obs_expon
autocorrelation_independency_unitroot_check(xt)



```


autocorrelation: [ 1.00000000e+00  3.88956586e-02 -1.27107166e-02 -1.87121182e-
 2.57412476e-02  4.87510993e-03 -7.17329648e-03 -3.25925063e-02
 6.82812566e-04  5.80513666e-04 -1.61784799e-02]
Ljung Box q statistic: [ 3.03028541  3.35405694  3.36107734  4.69028463  4.73798481
 6.97544702  6.97638417  6.97706189  7.50370501]
one-tailed p-value: [0.081723639788411018, 0.18692861571470631, 0.33923188257387138

```

```

D:\Softwares\Anaconda3\envs\py36\lib\site-packages\ipykernel\__main__.py:5: FutureV
D:\Softwares\Anaconda3\envs\py36\lib\site-packages\ipykernel\__main__.py:6: FutureV

```




```

Test Statistic          -42.975524
p-value                  0.000000
#Lags Used               0.000000
Number of Observations Used 1999.000000
Critical Value (1%)      -3.433625
Critical Value (5%)      -2.862987
Critical Value (10%)     -2.567540
dtype: float64

```


xt is independent and do not has a unit root

```
In [179]: yt0 = gen_hybrid_time_series(xt,0)
          autocorrelation_independency_unitroot_check(yt0)
```



```
autocorrelation: [ 1.00000000e+00  3.69241772e-01  2.83635539e-05  2.50595208e-
 2.58701560e-02  1.07381447e-02 -1.63208196e-02 -3.37909812e-02
-1.01821937e-02 -4.76353758e-03 -3.94713525e-03]
Ljung Box q statistic: [ 272.95185658  272.95185819  272.96444295  274.30632726  27
 275.07224594  277.36507752  277.57336892  277.61897954  277.65031161]
one-tailed p-value: [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0]
```

```
D:\Softwares\Anaconda3\envs\py36\lib\site-packages\ipykernel\__main__.py:5: FutureW
D:\Softwares\Anaconda3\envs\py36\lib\site-packages\ipykernel\__main__.py:6: FutureW
```



```


Test Statistic          -22.818631
p-value                  0.000000
#Lags Used               2.000000
Number of Observations Used 1996.000000
Critical Value (1%)      -3.433630
Critical Value (5%)      -2.862989
Critical Value (10%)     -2.567541
dtype: float64

```

```


In [181]: yt1 = gen_hybrid_time_series(xt,0.1)
          autocorrelation_independency_unitroot_check(yt1)

```

```
autocorrelation: [ 1.00000000e+00  3.68995012e-01  1.71510421e-05  2.13336122e-
 2.56564079e-02  1.07556806e-02 -1.70859869e-02 -3.42354618e-02
-1.01360931e-02 -4.35399200e-03 -4.05091893e-03]
Ljung Box q statistic: [ 272.58715672  272.58715731  272.59627801  273.91607972  27
 274.73405763  277.08760492  277.29401448  277.33211949  277.36512088]
one-tailed p-value: [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0]
```

```
D:\Softwares\Anaconda3\envs\py36\lib\site-packages\ipykernel\__main__.py:5: FutureW
D:\Softwares\Anaconda3\envs\py36\lib\site-packages\ipykernel\__main__.py:6: FutureW
```



```


Test Statistic          -22.831539
p-value                  0.000000
#Lags Used               2.000000
Number of Observations Used 1996.000000
Critical Value (1%)      -3.433630
Critical Value (5%)      -2.862989
Critical Value (10%)     -2.567541
dtype: float64

```

```


In [182]: yt2 = gen_hybrid_time_series(xt,1)
          autocorrelation_independency_unitroot_check(yt2)

```



```
autocorrelation: [ 1.00000000e+00  3.65142397e-01 -4.36598280e-04  3.97404489e-
 2.62877339e-02  9.79415079e-03 -1.63502247e-02 -3.50250856e-02
-1.66692234e-02 -7.65936814e-03 -3.97462878e-03]
Ljung Box q statistic: [ 266.92479847  266.92518028  266.95682964  268.34238309  26
 269.07134855  271.53471464  272.09295259  272.21087401  272.24264408]
one-tailed p-value: [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0]
```

```
D:\Softwares\Anaconda3\envs\py36\lib\site-packages\ipykernel\__main__.py:5: FutureW
D:\Softwares\Anaconda3\envs\py36\lib\site-packages\ipykernel\__main__.py:6: FutureW
```



```


Test Statistic          -22.831083
p-value                  0.000000
#Lags Used               2.000000
Number of Observations Used 1996.000000
Critical Value (1%)      -3.433630
Critical Value (5%)      -2.862989
Critical Value (10%)     -2.567541
dtype: float64

```

```


In [183]: yt3 = gen_hybrid_time_series(xt,10)
          autocorrelation_independency_unitroot_check(yt3)

```




```
autocorrelation: [ 1.          0.19232639  0.01121018  0.02926424  0.01118929 -0.00070000  
-0.00445739 -0.03327282 -0.02266939 -0.02224934  0.00646159]  
Ljung Box q statistic: [ 74.05291375  74.30462744  76.02085127  76.2718793   76.43610000  
76.47687047  78.69992303  79.73237141  80.72741248  80.81137856]  
one-tailed p-value: [0.0, 1.1102230246251565e-16, 2.2204460492503131e-16, 1.1102230246251565e-16,
```

```
D:\Softwares\Anaconda3\envs\py36\lib\site-packages\ipykernel\__main__.py:5: FutureWarning:
D:\Softwares\Anaconda3\envs\py36\lib\site-packages\ipykernel\__main__.py:6: FutureWarning:
```




Test Statistic	-36.615660
p-value	0.000000
#Lags Used	0.000000
Number of Observations Used	1998.000000
Critical Value (1%)	-3.433627
Critical Value (5%)	-2.862988
Critical Value (10%)	-2.567541
dtype:	float64

```
In [184]: yt4 = gen_hybrid_time_series(xt,100)
          autocorrelation_independency_unitroot_check(yt4)
```



```
autocorrelation: [ 1.          -0.00259711 -0.03625466 -0.04530756 -0.01008374 -0.01008374
 0.0200012  -0.00227503  0.04345297  0.04202723  0.03669171]
Ljung Box q statistic: [ 1.35034850e-02  2.64625232e+00  6.76002932e+00  6.96396372e+00
 7.67466372e+00  8.47756986e+00  8.48796293e+00  1.22813541e+01
 1.58316812e+01  1.85391348e+01]
one-tailed p-value: [0.092509526734801395, 0.26630149848911966, 0.07995278059424793, 0.07995278059424793,
0.07995278059424793, 0.07995278059424793, 0.07995278059424793, 0.07995278059424793,
0.07995278059424793, 0.07995278059424793]
```

```
D:\Softwares\Anaconda3\envs\py36\lib\site-packages\ipykernel\__main__.py:5: FutureWarning: The
D:\Softwares\Anaconda3\envs\py36\lib\site-packages\ipykernel\__main__.py:6: FutureWarning: The
```



```

Test Statistic          -44.811330
p-value                  0.000000
#Lags Used               0.000000
Number of Observations Used 1998.000000
Critical Value (1%)      -3.433627
Critical Value (5%)      -2.862988
Critical Value (10%)     -2.567541
dtype: float64

```

```

In [188]: from statsmodels.tsa.stattools import acf
          acorrs = [acf(yt0,nlags=1)[1], acf(yt1,nlags=1)[1], acf(yt2,nlags=1)[1],
                    acf(yt4,nlags=1)[1]]
          sigmas = [0,0.1,1,10,100]
          df = pd.DataFrame()
          df['sigma'] = sigmas
          df['1st_autocorr']=acorrs
          print(df)

```

```

      sigma  1st_autocorr
0      0.0      0.369242
1      0.1      0.368995
2      1.0      0.365142
3     10.0      0.192326

```



```
4 100.0 -0.002597
```

conclusion1: no matter how great the variance of white noise is, yt does not have a unitroot which indicates that it do not seems like a random walk.

conclusion2: The autocorrelation of 1 lag is positive with no white noise, and it decreases to 0 as far as variance of white noise increases.

1.2 Problem 2

Choose an equity with a long enough history and download the following data:

- Daily data for the past three years.
- Weekly data for the past 5 years.
- Monthly data for the past 10 years.
- Monthly data for the past 20 years (this will be used to calculate quarterly returns).

Calculate simple and continuously compounded returns for each dataset. Then construct four dataframes with 3 columns each (Price, SimpleReturn and LogReturn) and print the beginning of their content with the command head.

```
In [254]: import datetime
def preprocess_wrds_data(df):
    df['date'] = df['date'].apply(lambda x: datetime.datetime.strptime(str(x), '%Y-%m-%d'))
    # print(df.head())
    df.index = df.date
    df.drop(['date'], axis=1, inplace=True)
    df['simple_return'] = (df.PRC - df.PRC.shift(1)) / df.PRC.shift(1)
    df['log_return'] = np.log(df.PRC / df.PRC.shift(1))
    # spdf = df.totval
    # print(spdf.head())
```

```
In [255]: df_3y_d = pd.read_csv('NVDA_3yr_d.csv')
print(df_3y_d.head())
preprocess_wrds_data(df_3y_d)
print(df_3y_d.head())
```

	PERMNO	date	PRC	RETX	sprtrn		
0	86580	20150102	20.130	0.003990	-0.000340		
1	86580	20150105	19.790	-0.016890	-0.018278		
2	86580	20150106	19.190	-0.030318	-0.008893		
3	86580	20150107	19.135	-0.002866	0.011630		
4	86580	20150108	19.860	0.037889	0.017888		

	PERMNO	PRC	RETX	sprtrn	simple_return	log_return
date						
2015-01-02	86580	20.130	0.003990	-0.000340	NaN	NaN
2015-01-05	86580	19.790	-0.016890	-0.018278	-0.016890	-0.017034
2015-01-06	86580	19.190	-0.030318	-0.008893	-0.030318	-0.030787

2015-01-07	86580	19.135	-0.002866	0.011630	-0.002866	-0.002870
2015-01-08	86580	19.860	0.037889	0.017888	0.037889	0.037189

```
In [256]: df_5y_d = pd.read_csv("NVDA_5yr_d.csv")
preprocess_wrds_data(df_5y_d)
print(df_5y_d.head())
take_last = lambda array_like: array_like[-1]
df_5y_w = df_5y_d.resample("W").apply(lambda array_like: array_like[-1])
df_5y_w['simple_return']=(df_5y_w.PRC-df_5y_w.PRC.shift(1))/df_5y_w.PRC.s
df_5y_w['log_return']=np.log(df_5y_w.PRC / df_5y_w.PRC.shift(1))
print(df_5y_w.head())
```

	PERMNO	PRC	RETX	sprtrn	simple_return	log_return
date						
2013-01-02	86580	12.7200	0.037520	0.025403	NaN	NaN
2013-01-03	86580	12.7300	0.000786	-0.002086	0.000786	0.000786
2013-01-04	86580	13.1500	0.032993	0.004865	0.032993	0.032460
2013-01-07	86580	12.7700	-0.028897	-0.003123	-0.028897	-0.029323
2013-01-08	86580	12.4915	-0.021809	-0.003242	-0.021809	-0.022050
	PERMNO	PRC	RETX	sprtrn	simple_return	log_return
date						
2013-01-06	86580	13.15	0.032993	0.004865	NaN	NaN
2013-01-13	86580	12.21	-0.001635	-0.000048	-0.071483	-0.074166
2013-01-20	86580	12.17	-0.006531	0.003403	-0.003276	-0.003281
2013-01-27	86580	12.41	0.018048	0.005445	0.019721	0.019529
2013-02-03	86580	12.37	0.008972	0.010053	-0.003223	-0.003228

```
In [257]: df_10y_m = pd.read_csv("NVDA_10yr_m.csv")
print(df_10y_m.head())
preprocess_wrds_data(df_10y_m)
print(df_10y_m.head())
```

	PERMNO	date	PRC	ALTPRC	ALTPRCDT		
0	86580	20070131	30.650	30.650	20070131		
1	86580	20070228	31.000	31.000	20070228		
2	86580	20070330	28.780	28.780	20070330		
3	86580	20070430	32.890	32.890	20070430		
4	86580	20070531	34.639	34.639	20070531		
	PERMNO	PRC	ALTPRC	ALTPRCDT	simple_return	log_return	
date							
2007-01-31	86580	30.650	30.650	20070131	NaN	NaN	
2007-02-28	86580	31.000	31.000	20070228	0.011419	0.011355	
2007-03-30	86580	28.780	28.780	20070330	-0.071613	-0.074307	
2007-04-30	86580	32.890	32.890	20070430	0.142808	0.133488	
2007-05-31	86580	34.639	34.639	20070531	0.053177	0.051812	

```
In [258]: df_20y_m = pd.read_csv("NVDA_20yr_m.csv")
          print(df_20y_m.head())
          preprocess_wrds_data(df_20y_m)
          print(df_20y_m.head())
```

	PERMNO	date	PRC	ALTPRC	ALTPRCDT		
0	86580	19981231	NaN	19.6875	19990122		
1	86580	19990129	19.0000	19.0000	19990129		
2	86580	19990226	21.9375	21.9375	19990226		
3	86580	19990331	21.1250	21.1250	19990331		
4	86580	19990430	18.2500	18.2500	19990430		

	PERMNO	PRC	ALTPRC	ALTPRCDT	simple_return	log_return
date						
1998-12-31	86580	NaN	19.6875	19990122	NaN	NaN
1999-01-29	86580	19.0000	19.0000	19990129	NaN	NaN
1999-02-26	86580	21.9375	21.9375	19990226	0.154605	0.143759
1999-03-31	86580	21.1250	21.1250	19990331	-0.037037	-0.037740
1999-04-30	86580	18.2500	18.2500	19990430	-0.136095	-0.146292

For each of the eight return time series, perform the following procedures:

- Calculate the ACF for each time series. Turn in the four plots for each return type side by side. Label each plot accordingly.

```
In [276]: from statsmodels.tsa.stattools import acf
          import statsmodels.api as sm

          acf3ys = acf(df_3y_d.simple_return.dropna(axis=0,inplace=False))
          acf3yl = acf(df_3y_d.log_return.dropna(axis=0,inplace=False))
          acf5ys = acf(df_5y_w.simple_return.dropna(axis=0,inplace=False))
          acf5yl = acf(df_5y_w.log_return.dropna(axis=0,inplace=False))
          acf10ys = acf(df_10y_m.simple_return.dropna(axis=0,inplace=False))
          acf10yl = acf(df_10y_m.log_return.dropna(axis=0,inplace=False))
          acf20ys = acf(df_20y_m.simple_return.dropna(axis=0,inplace=False))
          acf20yl = acf(df_20y_m.log_return.dropna(axis=0,inplace=False))
          fig = plt.figure(figsize=(20,20))
          ax1 = fig.add_subplot(421)
          fig = sm.graphics.tsa.plot_acf(df_3y_d.simple_return.dropna(axis=0,inplace=False),
                                         title="autocorr of simple return of 3 years da
          ax2 = fig.add_subplot(422)
          fig = sm.graphics.tsa.plot_acf(df_3y_d.log_return.dropna(axis=0,inplace=False),
                                         title="autocorr of log return of 3 years da
          ax1 = fig.add_subplot(423)
          fig = sm.graphics.tsa.plot_acf(df_5y_w.simple_return.dropna(axis=0,inplace=False),
                                         title="autocorr of simple return of 5 years
          ax2 = fig.add_subplot(424)
          fig = sm.graphics.tsa.plot_acf(df_5y_w.log_return.dropna(axis=0,inplace=False),
                                         title="autocorr of log return of 5 years we
```

```

ax1 = fig.add_subplot(425)
fig = sm.graphics.tsa.plot_acf(df_10y_m.simple_return.dropna(axis=0,inplace=True)
                               title="autocorr of simple return of 10 years monthly")

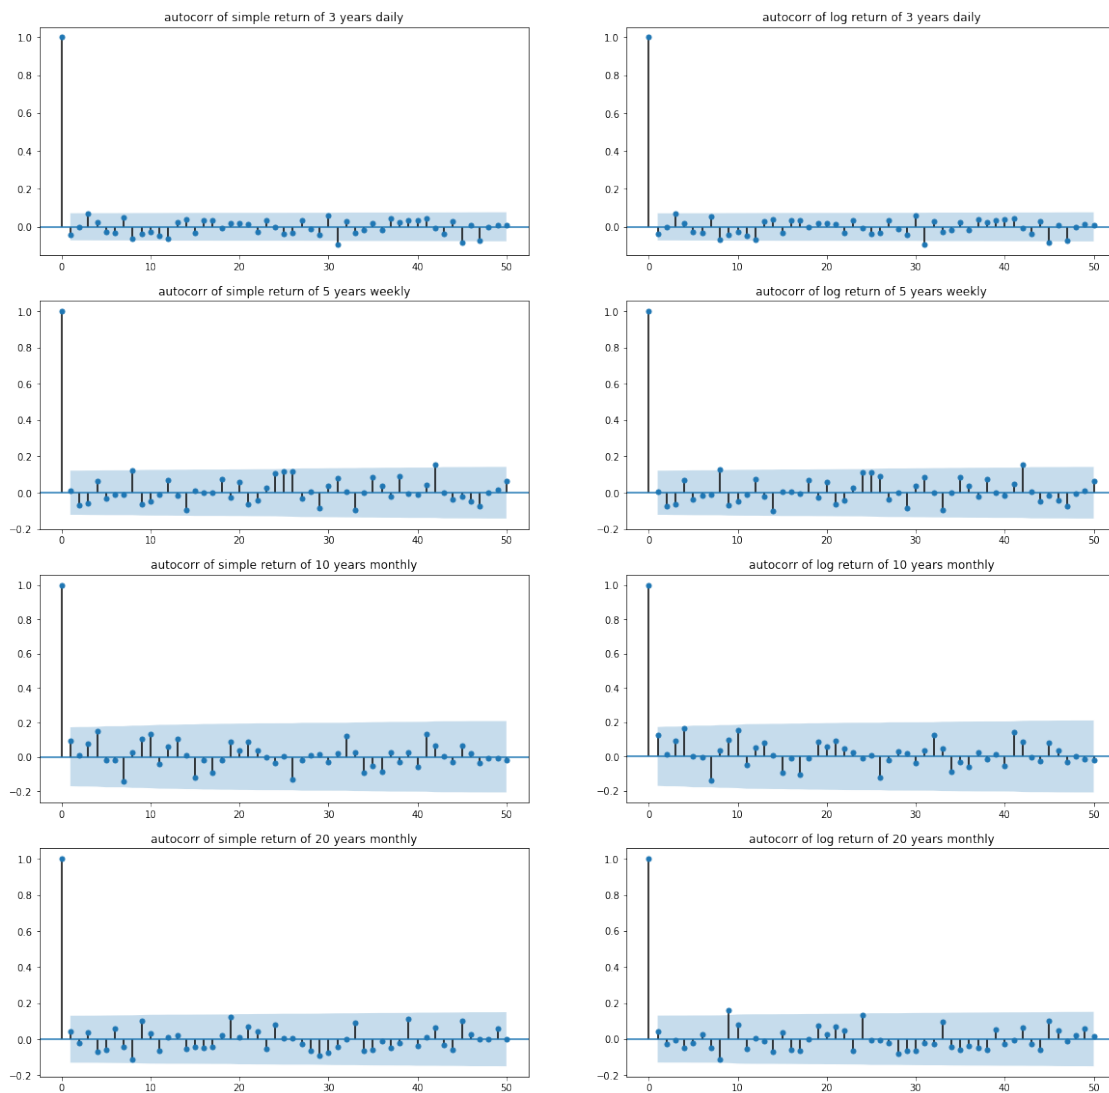
ax2 = fig.add_subplot(426)
fig = sm.graphics.tsa.plot_acf(df_10y_m.log_return.dropna(axis=0,inplace=True)
                               title="autocorr of log return of 10 years monthly")

ax1 = fig.add_subplot(427)
fig = sm.graphics.tsa.plot_acf(df_20y_m.simple_return.dropna(axis=0,inplace=True)
                               title="autocorr of simple return of 20 years monthly")

ax2 = fig.add_subplot(428)
fig = sm.graphics.tsa.plot_acf(df_20y_m.log_return.dropna(axis=0,inplace=True)
                               title="autocorr of log return of 20 years monthly")

plt.show()

```



(b) Perform the Portmanteau Ljung&Box test with as much lag as you feel is necessary.

```

In [284]: ljung_box_q_test(df_3y_d.simple_return.dropna(axis=0,inplace=False),m=10)
ljung_box_q_test(df_3y_d.log_return.dropna(axis=0,inplace=False),m=10)
print("*****")
ljung_box_q_test(df_5y_w.simple_return.dropna(axis=0,inplace=False),m=10)
ljung_box_q_test(df_5y_w.log_return.dropna(axis=0,inplace=False),m=10)
print("*****")
ljung_box_q_test(df_10y_m.simple_return.dropna(axis=0,inplace=False),m=10)
ljung_box_q_test(df_10y_m.log_return.dropna(axis=0,inplace=False),m=10)
print("*****")
ljung_box_q_test(df_20y_m.simple_return.dropna(axis=0,inplace=False),m=10)
ljung_box_q_test(df_20y_m.log_return.dropna(axis=0,inplace=False),m=10)

autocorrelation: [ 1.          -0.03990552 -0.0032686   0.06894479  0.02250864 -0.02994263
-0.02994263  0.0507622  -0.06216868 -0.03859024 -0.02612522]
Ljung Box q statistic: [ 1.2054914   1.21358979  4.82150374  5.20656549  5.781646432031
6.46432031  8.43063666 11.38386859 12.52330926 13.04623517]
one-tailed p-value: [0.27222688109693816, 0.4549048410540012, 0.18534393465301524,
autocorrelation: [ 1.          -0.03758476 -0.00293974  0.0682893   0.0194679  -0.03025393
-0.03025393  0.05397077 -0.06526339 -0.03986245 -0.02577275]
Ljung Box q statistic: [ 1.06935469  1.07590546  4.615541   4.9035924   5.486618384804
6.18384804  8.4065934 11.6611627 12.8769703 13.38588133]
one-tailed p-value: [0.30109125140929838, 0.41605748331579551, 0.2022129430140841,
*****
autocorrelation: [ 1.          0.0092385  -0.07111149 -0.06030858  0.0639766  -0.00846034
-0.00846034 -0.01132605  0.12440832 -0.06535443 -0.0471243 ]
Ljung Box q statistic: [ 0.02244803  1.35761174  2.32166208  3.41078527  3.69999806
3.75373344  7.93755331  9.09673375  9.70183008]
one-tailed p-value: [0.11909862962639972, 0.49277768093075558, 0.49161626564651772,
autocorrelation: [ 1.          0.00679744 -0.07570741 -0.06255142  0.0718728  -0.015162
-0.015162  -0.01004904  0.1272489  -0.0705048  -0.04755303]
Ljung Box q statistic: [ 0.01215249  1.5254758   2.56256436  3.93712471  4.303436538977
4.36538977  4.39257939  8.76963631 10.11871798 10.73487435]
one-tailed p-value: [0.087779645134105577, 0.46638775339002703, 0.46409032671815542,
*****
autocorrelation: [ 1.          0.09563432  0.01076806  0.07849631  0.14696467 -0.02214041
-0.02214041 -0.14182777  0.02763487  0.10533039  0.13210189]
Ljung Box q statistic: [ 1.22576479  1.24142539  2.08013474  5.04322664  5.103517175452
5.17175452  7.99809049  8.10626692  9.69068758 12.20347414]
one-tailed p-value: [0.26823226860591454, 0.4624388172588072, 0.44405791706653264,
autocorrelation: [ 1.          0.12238784  0.0111482   0.09255899  0.16219325  0.00289418
-0.00289418 -0.1372179  0.03581878  0.0990548   0.15440715]
Ljung Box q statistic: [ 2.00750275  2.02428859  3.19042709  6.79940941  6.799680085554
6.80085554  9.44644707  9.62818238 11.0294276 14.46241631]
one-tailed p-value: [0.15652278116842744, 0.36343882396750538, 0.36318659501706174,
*****
autocorrelation: [ 1.          0.04039932 -0.02236237  0.0377977  -0.07110838 -0.0597703
0.0597703  -0.04454189 -0.11072965  0.10359468  0.02974935]
Ljung Box q statistic: [ 0.37540591  0.4909412   0.82248749  2.00117343  2.790932

```

```

4.10003004  7.01037907  9.56943474  9.78144491]
one-tailed p-value: [0.45992777316883771, 0.21766371695718034, 0.15591889494414679,
autocorrelation: [ 1.          0.04072569 -0.02700104 -0.00396458 -0.05122357 -0.02
0.02551326 -0.04770646 -0.1149569  0.15947729  0.08223573]
Ljung Box q statistic: [ 0.38149591  0.54993387  0.55358148  1.16522179  1.300
1.45371082  1.99147635  5.1282799  11.19288501  12.81291383]
one-tailed p-value: [0.46319626649508372, 0.24040276131216654, 0.09302840287426184,

```

Conclusion: All return series are supposed to be independent under 1% significance level

- (c) For the famous Capital Asset Pricing Model (CAPM) to be valid, there should be no serial correlation between successive returns (data should be uncorrelated). Does this assumption appear to be valid based on your dataset? Compare the validity of this assumption using the various time frequencies.

From (a) and (b), we know the autocorrelations and Ljung Box p-value. Based on these results, roughly speaking, CAPM is correct. However, one can see the absolute value of autocorrelations are becoming higher when we increase the period of sampling, which conflicts with the iid assumption. Just note that, autocorrelation itself may not be independent. So we need more sophisticated model to capture the dynamics of it.

Then test the normality of each of the time series:

- (d) Do a normal qqplot for each of the time series. Turn in the four plots for each return type side by side. Label each plot accordingly.
- (e) Comment on the plots.

```

In [296]: fig = plt.figure(figsize=(10,20))
          ax1 = fig.add_subplot(421)
          fig = sm.qqplot(df_3y_d.simple_return.dropna(axis=0,inplace=False),ax=ax1)
          ax1.set_title("simple return of 3 years daily")

          ax2 = fig.add_subplot(422)
          fig = sm.qqplot(df_3y_d.log_return.dropna(axis=0,inplace=False),ax=ax2)

          ax2.set_title("log return of 3 years daily")

          ax3 = fig.add_subplot(423)
          fig = sm.qqplot(df_5y_w.simple_return.dropna(axis=0,inplace=False),ax=ax3)
          ax3.set_title("simple return of 5 years weekly")

          ax4 = fig.add_subplot(424)
          fig = sm.qqplot(df_5y_w.log_return.dropna(axis=0,inplace=False),ax=ax4)
          ax4.set_title("log return of 5 years weekly")


          ax5 = fig.add_subplot(425)
          fig = sm.qqplot(df_10y_m.simple_return.dropna(axis=0,inplace=False),ax=ax5)
          ax5.set_title("simple return of 10 years monthly")

```

```
ax6 = fig.add_subplot(426)
fig = sm.qqplot(df_10y_m.log_return.dropna(axis=0,inplace=False),ax=ax6)
ax6.set_title("log return of 10 years monthly")

ax7 = fig.add_subplot(427)
fig = sm.qqplot(df_20y_m.simple_return.dropna(axis=0,inplace=False),ax=ax7)
ax7.set_title("simple return of 20 years monthly")

ax8 = fig.add_subplot(428)
fig = sm.qqplot(df_20y_m.log_return.dropna(axis=0,inplace=False),ax=ax8)
ax8.set_title("log return of 20 years monthly")
plt.show()
```



(f) Test if return distributions are skewed. Summarize and comment.

From the q-q plot, the daily and weekly returns are both right skewed. The monthly return of 20 years data are left skewed. Others are less clear, as a result, we quantify the value.

```
In [300]: skew = []
          skew.append(sp.stats.skew(df_3y_d.simple_return.dropna(axis=0,inplace=False))
          skew.append(sp.stats.skew(df_3y_d.log_return.dropna(axis=0,inplace=False))
          skew.append(sp.stats.skew(df_5y_w.simple_return.dropna(axis=0,inplace=False))
          skew.append(sp.stats.skew(df_5y_w.log_return.dropna(axis=0,inplace=False))
          skew.append(sp.stats.skew(df_10y_m.simple_return.dropna(axis=0,inplace=False))
          skew.append(sp.stats.skew(df_10y_m.log_return.dropna(axis=0,inplace=False))
          skew.append(sp.stats.skew(df_20y_m.simple_return.dropna(axis=0,inplace=False))
          skew.append(sp.stats.skew(df_20y_m.log_return.dropna(axis=0,inplace=False))
          print("skewness:", skew)
```

```
skewness: [2.7995748447758304, 2.078745634552496, 1.319815386064654, 0.922326970754
```

(g) Test if return distributions are leptokurtic. Summarize and comment.

Surly all the return series are leptokurtic (fat-tailed) from q-q plots. We quantify it as follows:

```
In [302]: kurtosis = []
          kurtosis.append(sp.stats.kurtosis(df_3y_d.simple_return.dropna(axis=0,inplace=False))
          kurtosis.append(sp.stats.kurtosis(df_3y_d.log_return.dropna(axis=0,inplace=False))
          kurtosis.append(sp.stats.kurtosis(df_5y_w.simple_return.dropna(axis=0,inplace=False))
          kurtosis.append(sp.stats.kurtosis(df_5y_w.log_return.dropna(axis=0,inplace=False))
          kurtosis.append(sp.stats.kurtosis(df_10y_m.simple_return.dropna(axis=0,inplace=False))
          kurtosis.append(sp.stats.kurtosis(df_10y_m.log_return.dropna(axis=0,inplace=False))
          kurtosis.append(sp.stats.kurtosis(df_20y_m.simple_return.dropna(axis=0,inplace=False))
          kurtosis.append(sp.stats.kurtosis(df_20y_m.log_return.dropna(axis=0,inplace=False))
          print("kurtosis:", kurtosis)
```

```
kurtosis: [27.856972646838468, 19.82654671497985, 5.919559388482298, 4.071088839878
```

(h) Finally, test for normality of the returns. Perform a Jarque-Bera test as well as a Shapiro-Wilks test to check normality for each of the data under consideration. What do you conclude? Explain.

```
In [303]: def normality_test(array_like):
          jbvalue, jbp = sp.stats.jarque_bera(array_like)
          swvalue, swp = sp.stats.shapiro(array_like)
          print("Jarque Bera statistic:", jbvalue)
          print("Jarque Bera p-value:", jbp)
          print("Shapiro statistic:", swvalue)
          print("Shapiro p-value:", swp)
```

```
In [306]: normality_test(df_3y_d.simple_return.dropna(axis=0,inplace=False))
normality_test(df_3y_d.log_return.dropna(axis=0,inplace=False))
print("*****")
normality_test(df_5y_w.simple_return.dropna(axis=0,inplace=False))
normality_test(df_5y_w.log_return.dropna(axis=0,inplace=False))
print("*****")
normality_test(df_10y_m.simple_return.dropna(axis=0,inplace=False))
normality_test(df_10y_m.log_return.dropna(axis=0,inplace=False))
print("*****")
normality_test(df_20y_m.simple_return.dropna(axis=0,inplace=False))
normality_test(df_20y_m.log_return.dropna(axis=0,inplace=False))
```

```
Jarque Bera statistic: 25364.6040554
Jarque Bera p-value: 0.0
Shapiro statistic: 0.8134927153587341
Shapiro p-value: 1.2294258059430987e-28
Jarque Bera statistic: 12892.6676239
Jarque Bera p-value: 0.0
Shapiro statistic: 0.8454837799072266
Shapiro p-value: 1.8616553450928162e-26
*****
Jarque Bera statistic: 455.095701309
Jarque Bera p-value: 0.0
Shapiro statistic: 0.9199478626251221
Shapiro p-value: 1.3320058944721325e-10
Jarque Bera statistic: 216.412218816
Jarque Bera p-value: 0.0
Shapiro statistic: 0.9436067938804626
Shapiro p-value: 1.893412715503473e-08
*****
Jarque Bera statistic: 26.8055559615
Jarque Bera p-value: 1.51094091194e-06
Shapiro statistic: 0.9673646688461304
Shapiro p-value: 0.0030423414427787066
Jarque Bera statistic: 22.5024078277
Jarque Bera p-value: 1.29916474108e-05
Shapiro statistic: 0.9671876430511475
Shapiro p-value: 0.002931358991190791
*****
Jarque Bera statistic: 56.7026414478
Jarque Bera p-value: 4.86610751693e-13
Shapiro statistic: 0.9618170857429504
Shapiro p-value: 9.150791811407544e-06
Jarque Bera statistic: 268.363761206
Jarque Bera p-value: 0.0
Shapiro statistic: 0.9320725202560425
Shapiro p-value: 9.629804509359019e-09
```

None of them pass the test. The return series at any frequency are not normal. At least, we should not consider return as a Gaussian random variable. That is why we introduce Levy process family(e.g. generalized hyperbolic model, Merton jump, Kou) and stochastic volatility models(e.g. Heston, Heston-Nandi, GARCH family)

In []: